

TAI Me Up!

A Detailed Look at the Transponder
Abstraction Interface

Scott Emery

Terminology

- Optical transceiver
 - A hardware component which converts electrical signals to/from light
 - Sometimes called a transponder
 - May or may not adhere to some form of the MSA MIS
- TAI Adapter
 - The hardware-specific component of TAI
 - A user-mode device driver
 - Implemented as a shared object library: libtai.so
 - Sometimes called a driver or device-specific code
 - Not thread-safe, non-reentrant
- TAI Adapter Host
 - The code which interfaces to the TAI adapter
 - Sometimes called a TAI host, application, or device-independent code
 - Typically a daemon or systemd service

Terminology (cont.)

- TAI Module

- The combination of a DSP and optical transceiver
 - CFP2 DCO modules
 - CFP2 ACO modules plus the external DSP
 - An AC400 5x7 module

- Network Interface

- An optical interface, which connects externally to some type of optical line system
- The “external” optical port(s) of an optical module

- Host Interface

- The datapath electrical interface(s) of the TAI module
- In Voyager and Cassini are connected to the Broadcom Tomahawk ASIC

Terminology (cont.)

- TAI Object

- Currently one of these three types: Module, Host Interface, Network Interface
 - Also called a TAI API
 - Each object type has a set of attributes
- Created for each instance in the system, e.g. a network interface object is created for each network interface in the system
 - Each object instance is given an object ID

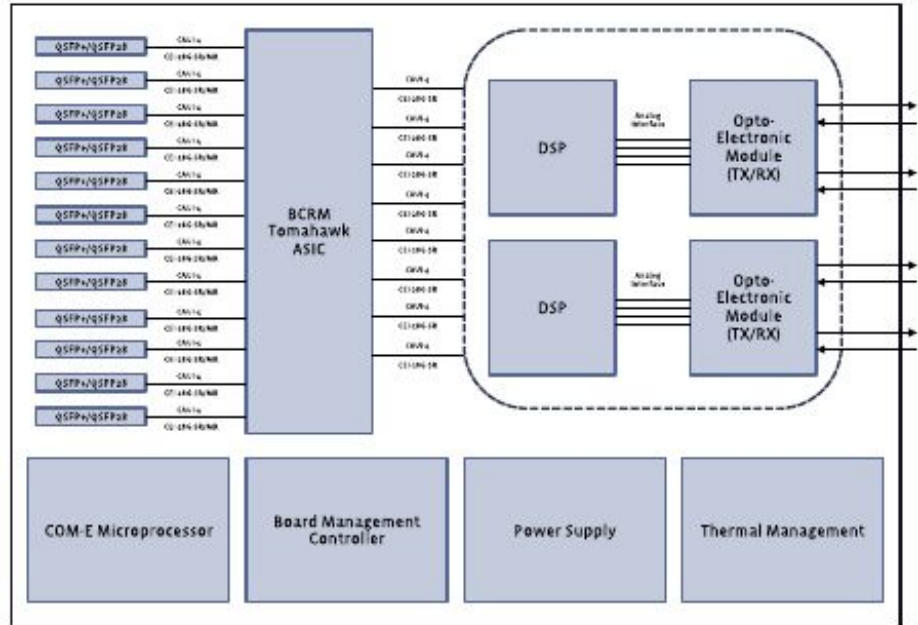
- TAI Attribute

- A property of an object instance which can be read and possibly written
- For example:
 - Module objects have a serial number attribute which can be read
 - Network interface objects have an output power attribute which can be read/written
- Types are defined in `taitypes.h`
- Object attributes defined in enums in `taimodule.h`, `tainetworkif.h`, and `tailhostif.h`
- Custom attributes are allowed

Voyager High Level



- Tomahawk
 - 12 “Client” ports (swpX, QSFP28)
 - 8 “Host” ports to AC400s (swpLX)
- COM-e Microprocessor
 - 4-core Atom E3845 @ 1.9GHz
 - 8G DRAM
 - 128G SSD
- BMC
 - AST2520
 - OpenBMC (not upstreamed)
 - Random eth0 MAC address
- 2 x AC400
 - Acacia
 - 2 x 200G DWDM each




Cassini High Level

- Tomahawk+
 - 16 “Host” ports (QSFP28)
 - 8 module slots (2 CAUI-4 each)
- CPU Subsystem
 - 4-core Broadwell-DE 1.6GHz
 - 8G DRAM
 - 32G Flash
- Module slots accept different cards
 - LCDCO-1
 - CFP2 DCO module slot
 - PHY for MACsec
 - LCQSFP-1
 - 2 QSFP28 modules
 - PHY for MACsec
 - LCACO-1
 - CFP2 ACO module slot
 - DSP




The TAI Interface

- Only 7 public symbols are exposed by TAI adapter
 - tai_api_initialize(), tai_api_query(), tai_api_uninitialize()
 - tai_log_set(), tai_object_type_query(), tai_module_id_query(), tai_dbg_generate_dump()
- Prototypes defines in tai.h
- But... several of these functions pass additional function pointers




tai_api_initialize()

- Provides initial entry point for TAI adapter code
 - Initialize variables
 - Allocate memory
 - Start threads/hook ISRs
- Does NOT touch the hardware
- Receives flags (unused) and pointer to a Service table (function pointers in TAI host)
 - Currently only one function pointer: module_presence()
- Not intended to take a long time, should quickly initialize and return




module_presence()

- Is called whenever a module is detected or removed
- Can be called in any context
 - Different thread
 - ISR
 - Before tai_api_initialize() returns
- Since TAI adapter code is not reentrant care must be taken to not call TAI adapter recursively
- Passes a presence/absence flag and module location (string)
 - Module location string is used in subsequent module object creation function call



tai_api_query()

- Retrieve a method table (function entry points) for an object type
- Allows TAI host to get the TAI adapter methods for an object type
 - Module
 - Host interface
 - Network Interface
- Returns a pointer to a structure in TAI adapter with function entry points
 - Structure must remain valid until tai_api_uninitialize() is called



tai_api_query() - cont.


- Currently the methods of all objects are similar:

- `create_module()`
- `remove_module()`
- `set_module_attribute()`
- `set_module_attributes()`
- `get_module_attribute()`
- `get_module_attributes()`
 - Replace “module” with “network_interface” or “host_interface”

- But this may change with additional capabilities


- Statistics, Performance monitoring, FAWS, etc.

- More on these functions later




tai_api_uninitialize()

- Provides for TAI adapter code to clean up
 - Free memory
 - Stop threads/un-hook ISRs
- No parameters
- Typically called when TAI Host exits




tai_log_set()

- Sets the logging level in the TAI adapter
- Two parameters
 - TAI API (object type)
 - Different logging levels can be set for each object type
 - Logging level
 - Follows standard syslog logging levels
 - Debug, Info, Notice, Warn, Error, Critical
- Enables logging above and including the selected level
 - E.g. Selecting Warn enables Warn, Error, and Critical messages
- Default logging level is Warn




tai_object_type_query()

- All object IDs are 64-bit opaque values provided by the TAI Adapter to the TAI Host
 - Each TAI Adapter defines the format of the object IDs used by that TAI Adapter
- Given an object ID, return the object type
- For example, passing in an object ID for a network interface object, will return that the object is a network interface type




tai_module_id_query()

- Given any object ID, return the module ID to which that object belongs
- For example, network interface and host interface objects “belong” to a module. Calling this function with a network interface object ID will return the module object ID to which that network interface belongs.
- Calling with a module object ID will simply return that same module object ID




tai_dbg_generate_dump()

- Requests that the TAI adapter dump its state to a file for debugging purposes
- A filename (char *) is passed
- The format of the dump file is not specified by TAI
 - Specific for each TAI adapter




create_module()

- Called to create a module instance
 - Typically invoked as a result of a module_presence() call with the present flag set
- TAI host supplies
 - A list of attributes to set
 - Must include the module location (value provided by the module_presence() function)
 - Pointer to a function notification table
 - Allows TAI adapter to notify TAI host
 - Shutdown request
 - State change
- TAI adapter
 - Initializes module to default state and applies attributes
 - Returns module object ID




shutdown_request()

- Called by the TAI Adapter to request the module be removed
 - An unrecoverable error
 - For maintenance
- TAI adapter supplies the module ID
- Will typically result in a `module_remove()`
 - Again, be careful to avoid re-entrant code




state_change_notification()

- Called by the TAI Adapter to provide notification of a change of state in the module
 - E.g. Initialize or Ready
- TAI adapter supplies the module ID and the new state
- Adapter host can react (or ignore) these changes



remove_module()

- Called to remove a module instance
 - Typically invoked as a result of a module_presence() call with the present flag clear, or when exiting
- TAI host supplies the module object ID
- TAI adapter
 - May or may not reset module to default state
 - Cleans up any resources allocated to the module
 - Question: Should it “remove” any network interface and host interface objects that are part of the module and not yet removed?




Module Attributes

- Most attributes for a module are read-only. Two are **read-write**
- Location - Somewhat opaque character string. Required on create, can only be supplied in module_create() attribute list
- Vendor name, part number, serial number, firmware versions
- Operational status, **administrative status**
- Temperature, input voltage
- Number of network interfaces, host interfaces
- **Tributary mapping**


set_module_attribute()

- Called to set the value of a single attribute
- TAI host supplies
 - module object ID
 - Pointer to an attribute
- TAI adapter
 - Performs whatever operations are necessary to modify the hardware to the supplied attribute value




set_module_attributes()

- Called to set the value of multiple attributes on a single module
- TAI host supplies
 - module object ID
 - The number of attributes in the list
 - Pointer to a list of attributes
- TAI adapter
 - Typically loops through the supplied attribute list calling set_module_attribute()




get_module_attribute()

- Called to retrieve the value of a single attribute
- TAI host supplies
 - module object ID
 - Pointer to an attribute with the ID to retrieve
- TAI adapter
 - Performs whatever operations are necessary to read the current value of the attribute
 - Put the retrieved value in the location provided




get_module_attributes()

- Called to retrieve the value of multiple attributes on a single module
- TAI host supplies
 - module object ID
 - The number of attributes in the list
 - Pointer to a list of attributes with attribute IDs
- TAI adapter
 - Typically loops through the supplied attribute list calling get_module_attribute()




create_network_interface()

- Called to create a network interface instance
 - NUM_NETWORK_INTERFACES attribute of module object can be used to determine how many network interface instances to create
- TAI host supplies
 - A list of attributes to set
 - Must include the network interface index (zero-based) on the module
 - The module ID upon which the network interface exists
- TAI adapter
 - Initializes network interface to default state and applies attributes
 - Returns network interface object ID
- NOTE: There are no notification functions



remove_network_interface()

- Called to remove a network interface instance
 - Typically invoked prior to calling remove_module()
- TAI host supplies the network interface object ID
- TAI adapter
 - May or may not reset network interface to default state
 - Cleans up any resources allocated to the network interface



Network Interface Attributes


- Many attributes are **read-write**, some are read-only
- **Index** - Zero-based index of network interfaces on a module. Required on create, can only be supplied in `network_interface_create()` attribute list
- **TX Enable** - This is inversion of the well-known TX_DIS
- **Grid Spacing, Output power**, Measured output power
- **TX Channel number, Channel by frequency, channel by lambda** - Overlap each other, three ways to set the same thing
- **Modulation format, Differential encoding**
- **Operational Status**
- **Pulse shaping: RX, TX, RX Beta, TX Beta**
- **RX VOA attenuation**

Network Interface Attributes (cont.)

- Current Laser Freq, Fine-tune laser freq, Min/Max Laser freq, Grid support
- TX/RX Alignment status
- BER, BER Period


set/get_network_interface_attribute/s()

- Same idea as the module object methods of the similar name




create_host_interface()

- Called to create a host interface instance
 - NUM_HOST_INTERFACES attribute of module object can be used to determine how many network interface instances to create
- TAI host supplies
 - A list of attributes to set
 - Must include the host interface index (zero-based) on the module
 - The module ID upon which the network interface exists
- TAI adapter
 - Initializes host interface to default state and applies attributes
 - Returns host interface object ID
- NOTE: There are no notification functions



remove_host_interface()

- Called to remove a host interface instance
 - Typically invoked prior to calling remove_module()
- TAI host supplies the host interface object ID
- TAI adapter
 - May or may not reset host interface to default state
 - Cleans up any resources allocated to the host interface




Host Interface Attributes

- Many attributes are **read-write**, some are read-only
- Index - Zero-based index of host interfaces on a module. Required on create, can only be supplied in `host_interface_create()` attribute list
- Lane Faults, TX Alignment status
- **FEC Type**
- There will typically be some custom attributes for SERDES configuraiton

set/get_network_interface_attribute/s()

- Same idea as the module/network interface object methods of the similar name





Thank you!

for staying awake

Visit us at cumulusnetworks.com or follow us [@cumulusnetworks](https://twitter.com/cumulusnetworks)

© 2017 Cumulus Networks. Cumulus Networks, the Cumulus Networks Logo, and Cumulus Linux are trademarks or registered trademarks of Cumulus Networks, Inc. or its affiliates in the U.S. and other countries. Other names may be trademarks of their respective owners. The registered trademark Linux® is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a world-wide basis.